

ChemRxivQuest: A Curated Chemistry Question-Answer Database Extracted from ChemRxiv Preprints

Mahmoud Amiri^{a,b}, Thomas Bocklitz^{a,b}

May 9, 2025

^a Leibniz Institute of Photonic Technology, Member of Research Alliance “Leibniz Health Technologies”, Jena, 07745, Germany

^b Institute of Physical Chemistry, Friedrich Schiller University Jena, 07743, Germany

Abstract

The rapid expansion of chemistry literature poses significant challenges for researchers seeking to efficiently access domain-specific knowledge. To support advancements in chemistry-focused natural language processing (NLP), we present ChemRxivQuest, a curated dataset of 970 high-quality question-answer (QA) pairs derived from 155 ChemRxiv¹ preprints across 17 subfields of chemistry. Each QA pair is explicitly linked to its source text segment to ensure traceability and contextual accuracy. ChemRxivQuest was constructed using an automated pipeline that combines optical character recognition (OCR), GPT-4o-based QA generation, and a fuzzy matching technique for answer verification. The dataset emphasizes conceptual, mechanistic, applied, and experimental questions, enabling applications in retrieval-based QA systems, search engine development, and fine-tuning of domain-adapted large language models. We analyze the dataset’s structure, coverage, and limitations, and outline future directions for expansion and expert validation. ChemRxivQuest provides a foundational resource for chemistry NLP research, education, and tool development.

Keywords: Chemistry Question Answering, Natural Language Processing, Chemistry NLP Dataset, Large Language Models, Dataset Curation, Scientific Text Mining.

¹<https://chemrxiv.org/>

1 Introduction

Chemistry is a diverse and rapidly evolving scientific discipline, witnessing an unprecedented pace of research publications. Each year, millions of chemistry-related articles, reviews, and preprints are published across platforms such as peer-reviewed journals and open-access repositories like ChemRxiv[9]. For example, ChemRxiv alone receives thousands of new submissions annually, reflecting the vibrant and fast-paced nature of chemical research. However, this growing body of literature presents significant challenges for chemists, researchers, and students seeking to efficiently locate, extract, and apply domain-specific knowledge.

Traditional literature search methods—relying primarily on keyword-based queries through tools such as Web of Science², Google Scholar³, and SciFinder⁴—often yield vast quantities of unstructured documents. These require labor-intensive manual review, lack semantic understanding, and rarely provide direct, contextualized answers to targeted chemical questions. Researchers are left to synthesize information themselves, limiting efficiency and accessibility, particularly for interdisciplinary or novice users.

Recent breakthroughs in NLP⁵, especially LLMs⁶ and RAG⁷, offer promising new approaches to address these challenges. LLMs such as GPT-3 and GPT-4 have demonstrated the ability to perform tasks ranging from text classification and question answering to autonomous experiment planning and synthesis optimization in chemistry and materials science contexts[1, 2, 5, 11]. Systems like Coscientist and ChemCrow exemplify how LLMs can be integrated with external tools and knowledge bases to automate experimental workflows, data extraction, and even discovery of novel compounds[1, 11]. These models significantly reduce the dependence on large annotated datasets or complex domain-specific architectures[2, 5], lowering the barrier to entry for non-experts and accelerating knowledge-intensive research. Despite these advances, chemistry-focused NLP continues to lag behind other fields, such as biomedicine[12]. This is due in part to the intrinsic complexity of chemical language—featuring diverse terminologies, symbolic notations, and reaction mechanisms—and a lack of high-quality, structured, and curated datasets explicitly designed for training and evaluating chemistry-specific LLMs. Additionally, recent studies highlight the risks of hallucination and factual inaccuracies in generative models, which can undermine their utility in high-stakes scientific domains[3, 7]. Approaches such as fine-tuning, entropy-based uncertainty estimation, and RAG have been proposed to mitigate these issues and improve answer reliability[3, 4, 19].

To address these critical gaps, we introduce ChemRxivQuest, a curated and systemat-

²<https://www.webofscience.com/>

³<https://scholar.google.com/>

⁴<https://sso.cas.org/>

⁵natural language processing

⁶large language models

⁷retrieval-augmented generation

ically constructed question–answer dataset for chemistry-specific NLP applications. Built directly from primary literature on ChemRxiv, ChemRxivQuest is designed to support robust question answering, improve retrieval precision, and serve as a fine-tuning or benchmarking resource for domain-specialized LLMs. Our aim is to catalyze advances in chemical information retrieval and automated understanding by providing a reliable, structured, and versatile dataset tailored to the unique demands of chemical language.

Research on scientific QA⁸ datasets and chemical NLP has laid an important foundation for the development of ChemRxivQuest. Notable datasets in this space include PubMedQA[8], BioASQ[17], SciQ[18], and the S2ORC corpus[10]. While these datasets have supported significant progress in biomedical and general scientific domains, they exhibit limitations when applied to the unique demands of chemistry-focused NLP.

PubMedQA focuses heavily on biomedical literature, offering question–answer pairs grounded in clinical contexts. Despite its strengths in medical applications, it does not capture the chemical specificity and breadth required for tasks in chemical informatics. Similarly, SciQ is constructed from educational assessment materials and is tailored toward general science education rather than research-grade chemical content, thus lacking the granularity needed for specialized QA systems. BioASQ, like PubMedQA, is centered around biomedical text mining and question answering, but its domain restriction excludes complex chemical representations. S2ORC, although extensive and multidisciplinary, remains loosely structured and lacks chemistry-targeted QA annotations, which hinders its utility for precise chemical information retrieval.

In the domain of chemical NLP, tools such as ChemDataExtractor[16] and OSCAR[6] have made significant advances in information extraction. ChemDataExtractor has been particularly effective in mining structured data such as molecular properties, synthesis conditions, and reaction outcomes, while OSCAR has improved chemical named entity recognition through rule-based and machine learning approaches. However, neither tool is designed to support structured QA tasks or generative modeling tailored to chemical discourse.

Despite the availability of these valuable resources and tools, a clear gap persists in the form of a curated, validated QA dataset that is specifically tailored to chemistry and directly derived from primary scientific literature. ChemRxivQuest addresses this unmet need by providing an explicitly structured, chemistry-specific QA dataset with broad topical coverage and high-quality annotations suitable for downstream applications in chemical NLP.

In addressing the limitations outlined above, ChemRxivQuest introduces a number of novel contributions to the field of chemical NLP. It represents the first large-scale, curated QA dataset specifically focused on chemical literature, sourced from ChemRxiv preprints. The dataset comprises 970 QA pairs, each validated and explicitly linked to

⁸question answering

its original textual context to ensure traceability and scientific integrity.

The dataset construction process is underpinned by a rigorous and reproducible methodological pipeline. This includes OCR⁹ for accurate text extraction from PDF documents, automated QA pair generation using state-of-the-art large language models (specifically GPT-based architectures), and a verification step employing fuzzy matching algorithms to ensure alignment between generated answers and their source contexts. This approach ensures that the resulting dataset is both accurate and reliable for use in high-stakes chemical NLP applications.

ChemRxivQuest also ensures comprehensive topical coverage across 17 subfields of chemistry, including organic, inorganic, physical, analytical, and materials chemistry. This diversity enhances its utility as a benchmark dataset and training resource for a wide array of tasks within the chemical sciences. Beyond its content, ChemRxivQuest is designed with practical applicability in mind. It can support the development of chemistry-specific retrieval-based QA systems, serve as a benchmark for evaluating domain-adapted NLP models, and aid in the fine-tuning of generative language models for scientific applications. Through these contributions, ChemRxivQuest aims to significantly advance the capabilities of NLP systems in chemistry, enabling more effective access to and understanding of the chemical literature.

2 Dataset Construction

Creating *ChemRxivQuest* involved a systematic and clearly structured pipeline comprising several integral stages: data source selection, comprehensive text extraction and preprocessing, generation of high-quality QA pairs using advanced language models, and rigorous answer verification and indexing. This section outlines each of these stages to ensure methodological transparency and reproducibility.

2.1 Data Source: ChemRxiv Preprints

⁹optical character recognition

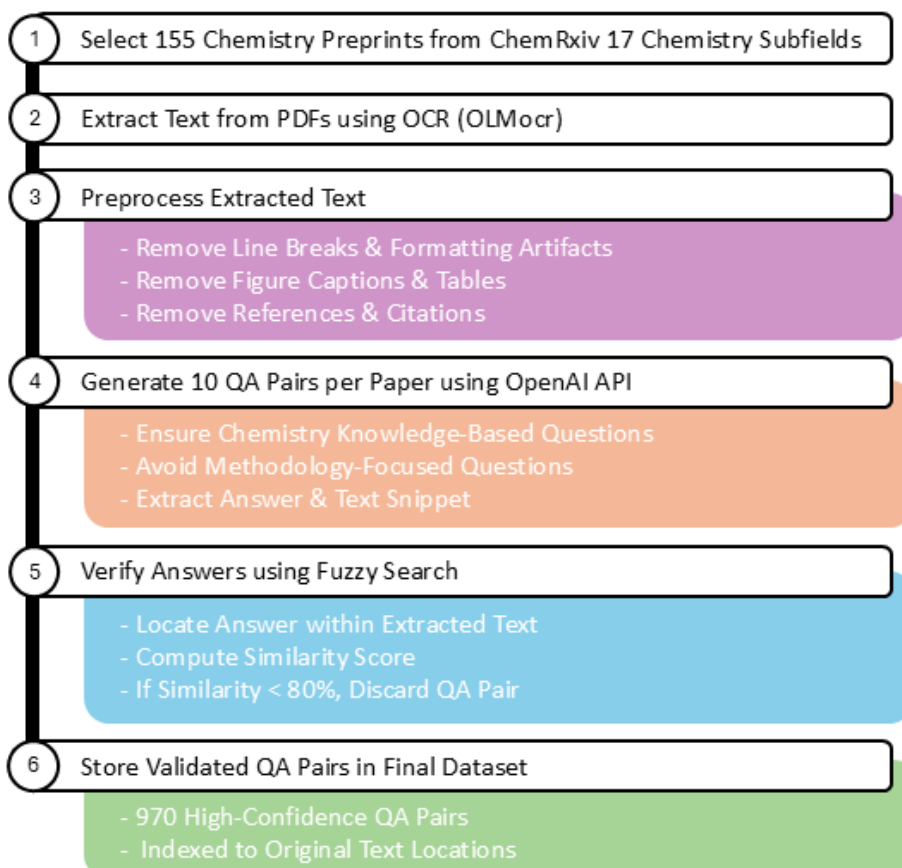


Figure 1: Overview of the ChemRxivQuest dataset construction pipeline. The process begins with the selection of 155 chemistry preprints from 17 ChemRxiv subfields, followed by text extraction using OCR. Extracted text is preprocessed to remove formatting artifacts, figures, and references. Using the OpenAI API, chemistry-focused question–answer pairs are generated and subsequently validated through fuzzy matching. Pairs with insufficient answer similarity are discarded. The final dataset consists of 970 high-confidence QA pairs, each indexed to its source text location.

We selected *ChemRxiv*, a prominent open-access repository for chemistry research, as the sole data source for ChemRxivQuest. *ChemRxiv* provides a broad collection of recent, high-quality preprints across many subfields of chemistry and operates under the Creative Commons CC BY¹⁰ license, which explicitly permits unrestricted reuse and adaptation—making it well-suited for constructing an open dataset.

To ensure disciplinary diversity and thematic balance, we curated a total of 155 preprints, evenly distributed across 17 well-established subfields of chemistry: *Agriculture and Food Chemistry, Analytical Chemistry, Biological and Medicinal Chemistry, Catalysis, Chemical Education, Chemical Engineering and Industrial Chemistry, Earth, Space, and Environmental Chemistry, Energy Chemistry, Inorganic Chemistry, Materials Chemistry, Materials Science, Nanoscience, Organic Chemistry, Organometallic Chemistry, Physical Chemistry, Polymer Science, and Theoretical and Computational Chemistry.*

Preprints were selected based on three criteria: recency (published within the past three years), subfield balance (ensuring roughly equal representation from each category), and quality (emphasizing clarity of presentation, depth of chemical content, and potential for generating meaningful conceptual or mechanistic questions). Metadata such as publication date, title, authors, and fields of study were automatically extracted using OpenAI’s GPT-3.5-turbo model and validated against a controlled vocabulary. This curated collection ensures a dataset that reflects both the breadth and depth of contemporary chemical research.

2.2 Text Extraction and Preprocessing

As most *ChemRxiv* preprints are distributed as PDFs, accurate text extraction was essential. To achieve this, we used *OLMocr*[14], an open-source OCR tool specifically optimized for scientific literature. This tool leverages CNN¹¹ trained on scientific texts, ensuring high fidelity in recognizing technical terminology and chemical nomenclature.

Each PDF was processed individually at 300 DPI resolution, converting scanned pages into machine-readable text. The output was then cleaned using Python scripts to remove OCR artifacts (e.g., misrecognized characters and redundant whitespace), along with extraneous content such as page numbers, headers, and footers.

Subsequently, we filtered out irrelevant content such as figure captions, tables, and reference sections. This was achieved using regex-based heuristics tailored to common formatting patterns in scientific writing. To ensure accuracy, a randomly selected 10% of documents were manually reviewed and used to iteratively refine the heuristics. The resulting preprocessed text was of sufficient quality for reliable question-answer generation.

¹⁰<https://creativecommons.org/licenses/by/4.0/>

¹¹convolutional neural networks

2.3 Question-Answer Generation Process

Question-answer QA pairs were generated using OpenAI’s GPT-4o model due to its strong performance on scientific reasoning tasks. The model was prompted with structured instructions to generate three specific types of questions: (1) conceptual questions explaining fundamental chemistry concepts, (2) mechanistic questions describing chemical processes, and (3) applied questions discussing the practical implications of the research.

Each prompt required the model to return a clearly stated question, a concise answer, and a directly relevant textual snippet quoted verbatim or nearly verbatim from the source preprint. This structure was designed to ensure that every QA pair was accurate, grounded in source material, and easy to validate.

To optimize output reliability, the GPT-4o model was configured with a temperature of 0.2 to minimize randomness. For each of the 155 preprints, 10 QA pairs were generated, yielding a total of 1,550 candidate pairs.

2.4 Answer Verification and Indexing

To ensure factual alignment and traceability, each generated answer was validated by checking its similarity to the original source text. A hybrid verification pipeline was implemented:

First, an exact substring match was attempted. If that failed, we applied approximate matching using the `fuzzysearch` and `rapidfuzz` libraries, leveraging Levenshtein-based[13] edit distance and partial ratio scores[15]. A match was accepted if its similarity exceeded an empirical threshold of 80%.

A sliding window approach (400-character windows, 100-character steps) was used to localize the best-matching passage. Matches were buffered by 10 characters on either side to capture partial truncations. QA pairs failing to meet the threshold were discarded.

A random 10% of the retained QA pairs were manually reviewed. This manual validation confirmed an estimated post-filtering accuracy exceeding 95%. From the original pool of 1,550 generated QA pairs, 970 passed the verification pipeline, yielding a 62.6% retention rate.

3 Dataset Analysis

This section presents a comprehensive quantitative and qualitative analysis of the *ChemRxivQuest* dataset. The dataset comprises 971 validated QA pairs, extracted from 155 *ChemRxiv* preprints. We analyze the distribution across chemical subfields, categorize the types of questions included to demonstrate the dataset’s diversity and utility.

3.1 Distribution Across Chemistry Subfields

To ensure broad representativeness, the dataset was stratified to provide substantial coverage across 17 major subfields of chemistry. As shown in Table 1, the distribution reflects the interdisciplinary nature of chemical research, where some QA pairs belong to multiple subfields.

Table 1: percentage of QA pairs per Chemistry Subfield in *ChemRxivQuest*

Chemistry Subfield	percentages
Materials Chemistry	19
Theoretical and Computational Chemistry	18
Inorganic Chemistry	17
Physical Chemistry	14
Organic Chemistry	10
Nanoscience	9
Biological and Medicinal Chemistry	9
Analytical Chemistry	8
Chemical Engineering & Industrial Chemistry	8
Catalysis	7
Energy Chemistry	7
Earth, Space, and Environmental Chemistry	7
Polymer Science	5
Materials Science	5
Agriculture and Food Chemistry	5
Organometallic Chemistry	5
Chemical Education	3

On average, each subfield contains approximately 92 QA pairs, with a standard deviation of around 47.6, indicating moderate variation in distribution. This stratification ensures that the dataset encompasses a wide range of chemical domains, from theoretical principles to applied materials science.

3.2 Analysis of Question Types

To investigate the diversity of the content more thoroughly, each QA pair in the dataset was classified into one of four primary categories, based on the conceptual scope and focus of the question. These categories include conceptual, mechanistic, applied, and synthetic or experimental questions. Conceptual questions are those that address fundamental principles or theoretical definitions, such as inquiries into the nature of chemical interactions—for example, “What defines a Lewis acid-base interaction?” Mechanistic questions, on the other hand, delve into the specifics of chemical processes or reaction pathways, exemplified by questions like “How does a palladium catalyst facilitate cross-coupling reactions?” Applied questions emphasize real-world implementations and technological

relevance, as seen in examples such as “What are potential applications of graphene oxide in energy storage?” Finally, synthetic and experimental questions focus on laboratory methods, reagents, and protocols, such as “Which reagent is commonly used for selective oxidation of alcohols?”

The distribution of these question types is summarized in Table 2. Conceptual questions constitute the largest portion of the dataset, accounting for 23.7% of all QA pairs, highlighting their central role in both chemistry education and the fine-tuning of large language models (LLMs). Mechanistic questions represent 24.9% of the dataset, followed by applied questions at 25.5%, illustrating the dataset’s strong alignment with both theoretical insight and practical utility. Although synthetic and experimental questions make up a similar fraction (25.7%), they provide essential support for real-world chemical workflows and enhance the dataset’s applicability to laboratory-oriented retrieval systems.

Table 2: Distribution of QA Pair Types in *ChemRxivQuest*

Question Type	Number of Questions	Percentage (%)
Conceptual	231	23.7
Mechanistic	242	24.9
Applied	248	25.5
Synthetic and Experimental	250	25.7
Total	971	100.0

Although numerical balance across subfields was intentionally maintained, some domains naturally yield more accessible or conceptually focused content, which may influence the type and complexity of the generated QA pairs. For instance, conceptual questions are more readily derived from educational or review-type texts, whereas experimental fields may pose challenges in question generation due to technical density. Future expansions of *ChemRxivQuest* may include expert-driven annotations to further assess and mitigate these biases. Incorporating difficulty ratings or ambiguity scores could enhance the dataset’s granularity and support more nuanced model evaluations.

4 Applications and Use Cases

ChemRxivQuest is designed as a versatile and robust resource for advancing a wide range of chemistry-specific natural language processing (NLP) applications. Its structured and validated question–answer format, along with explicit linkage to source preprints, enables numerous practical uses across research, development, education, and benchmarking. This section outlines key application areas where *ChemRxivQuest* provides direct value.

4.1 Retrieval-Based Chemistry Search Engines

ChemRxivQuest can serve as a valuable testbed for developing and evaluating chemistry-specific retrieval-based search systems. Given its grounding in real scientific literature, the dataset supports the development of domain-aware information retrieval frameworks that go beyond keyword-based matching. By leveraging the explicit question–context–answer structure, researchers can benchmark retrieval models (e.g., BM25, dense vector search with FAISS) in terms of recall@k, MRR, and contextual relevance. Furthermore, its integration with semantic retrieval systems using large language models (LLMs) or embedding-based methods enables evaluation of hybrid retrieval–generation pipelines in real-world chemistry scenarios.

4.2 Fine-Tuning Chemistry-Specific Large Language Models (LLMs)

Domain-specific fine-tuning is essential for improving the performance of large language models on specialized tasks where general pretraining falls short. In this context, *ChemRxivQuest* offers a high-quality corpus for supervised fine-tuning of chemistry-focused LLMs. Researchers aiming to create customized models—such as a domain-adapted variant of GPT (e.g., “ChemGPT”)—can utilize the dataset’s rich and diverse QA pairs to instill deeper domain knowledge into the model.

The fine-tuning process involves using chemistry-specific questions as input prompts and their corresponding validated answers as outputs. For instance, a prompt like “Explain the role of Lewis acids in Friedel–Crafts reactions” would be paired with an accurate, contextually relevant answer sourced from a ChemRxiv preprint. Training can be implemented through frameworks such as Hugging Face Transformers or the OpenAI fine-tuning API. Performance metrics such as BLEU, ROUGE, and exact match accuracy offer quantitative evaluation, while the explicit association of each QA pair with source snippets enables integration with advanced techniques like retrieval-augmented generation (RAG) and reinforcement learning from human feedback (RLHF).

4.3 Benchmarking NLP Models in Scientific Question Answering

Beyond its role as a training corpus, *ChemRxivQuest* functions as a rigorous benchmark for scientific question answering in the chemistry domain. Its validated and source-grounded QA pairs make it ideal for evaluating NLP models under standardized and reproducible conditions. Researchers can assess a variety of model architectures—including encoder–decoder transformers, retrieval-only models, and RAG-based hybrids—by measuring their performance on chemistry-specific tasks.

Standard evaluation procedures include dataset splits into training, validation, and test sets, with metrics such as Exact Match (EM), ROUGE-L, and BLEU providing

insight into both lexical and semantic accuracy. The domain richness of the dataset further allows for fine-grained evaluation by question type, sub-discipline (e.g., organic, inorganic, materials chemistry), or reasoning complexity. These capabilities support the systematic analysis of model strengths and weaknesses, informing the development of more robust and chemically literate NLP systems.

4.4 Evaluating Embedding Models for Scientific Text Retrieval

The dataset also enables systematic benchmarking of text embedding models in scientific and chemical contexts. Because each question in *ChemRxivQuest* is associated with a specific answer and its source document, researchers can assess how well various embedding models perform in retrieving semantically relevant contexts. Models such as SciBERT, Specter, E5, and OpenAI or Cohere embeddings can be evaluated using metrics like recall@k, nDCG, and semantic similarity to the gold answer.

This application is particularly valuable in designing high-performance retrieval layers in RAG systems or chemistry-aware semantic search tools. By comparing embedding models across a common evaluation framework, the dataset facilitates clearer understanding of which vector representations are best suited for chemical language.

4.5 Educational Applications in Chemistry

ChemRxivQuest also holds significant value for chemical education. Its structured QA format supports a variety of pedagogical tools, from interactive learning platforms and automated tutoring systems to quiz generation and self-assessment modules. Because each QA pair is validated and linked to a credible scientific source, it offers students and educators trustworthy material for exploring complex concepts, mechanisms, and applications in modern chemistry. Interactive systems can use these QA pairs to simulate question-driven inquiry or Socratic learning environments, encouraging deeper conceptual engagement.

4.6 Analyzing Knowledge Drift and Domain Adaptation in LLMs

Finally, *ChemRxivQuest* enables longitudinal studies on how large language models evolve over time in their knowledge of chemistry. By re-evaluating model outputs across different LLM versions (e.g., GPT-3.5 vs. GPT-4 vs. domain-finetuned variants), researchers can track changes in factual accuracy, terminology usage, and reasoning quality. This is particularly relevant in assessing knowledge drift—where model responses may shift as a function of pretraining data updates or architectural changes.

Additionally, the dataset provides a platform for measuring the impact of domain adaptation efforts. Models pre-trained on general corpora can be compared to models

fine-tuned on *ChemRxivQuest* to determine improvements in answer correctness, retrieval fidelity, and reasoning coherence in chemical domains.

5 Limitations

Despite its scope and utility, the *ChemRxivQuest* dataset exhibits several notable limitations. A primary concern stems from the reliance on large language models (LLMs) to automatically generate questions and answers. Although the generation process was carefully designed to minimize ambiguity and speculative reasoning, instances of model hallucination remain a challenge. In certain cases, LLM-generated answers referenced chemical reactions or mechanistic details not explicitly supported by the original preprint content, introducing plausible-sounding but incorrect information. These inaccuracies, if unchecked, could propagate misinformation and compromise the dataset’s scientific reliability. Moreover, the current answer verification approach—based on fuzzy token matching via Levenshtein distance with a fixed similarity threshold—prioritizes textual similarity over semantic correctness. As a result, syntactically accurate but contextually invalid answers may inadvertently be accepted into the final dataset.

Another significant limitation concerns inherent biases and imbalances within the dataset. Although *ChemRxivQuest* spans 17 chemistry subfields, the distribution is not uniform. Subfields such as Organic Chemistry and Materials Science are disproportionately represented, likely reflecting higher submission rates on ChemRxiv. In contrast, areas like Agricultural and Food Chemistry or Chemical Education remain underrepresented. This imbalance could introduce bias during model training, resulting in uneven performance and limiting the generalizability of the dataset across the broader spectrum of chemical disciplines.

A further limitation arises from the exclusive use of ChemRxiv preprints as source material. While preprints facilitate rapid and open dissemination of emerging research, they have not undergone formal peer review, raising questions about the reliability, completeness, and scientific rigor of the information they contain. Claims or mechanistic insights presented in these preprints may be revised, corrected, or even retracted upon peer review. Consequently, the reliance on preprint literature introduces an additional layer of uncertainty regarding the long-term validity and accuracy of QA pairs derived from such sources.

6 Future Work

To address current limitations and further enhance the quality and utility of the *ChemRxivQuest* dataset, several avenues for future development are envisioned. A central priority

involves incorporating human expert validation into the dataset construction process. By involving domain experts from various subfields of chemistry, a portion of the QA pairs can undergo rigorous manual review, allowing for in-depth accuracy assessment. Such validation would not only help quantify the frequency and nature of large language model (LLM) hallucinations but also provide critical feedback for refining low-quality content. Integrating human-in-the-loop mechanisms—where expert feedback is directly used to revise or filter questions and answers—could markedly improve the reliability and trustworthiness of the dataset.

Another key direction is addressing content imbalances across different chemistry subfields. Currently underrepresented areas such as Agricultural Chemistry, Environmental Chemistry, and Chemical Education may benefit from targeted curation strategies. Employing stratified sampling approaches based on subfield-specific quotas could help ensure that future dataset releases offer a more comprehensive and balanced representation of chemical knowledge. This intentional diversification would make *ChemRxivQuest* more inclusive and valuable for a wider range of applications.

In addition, expanding the dataset to include content from peer-reviewed literature represents a significant opportunity to enhance quality and credibility. Future work will focus on integrating material from reputable open-access sources such as PubChem, Europe PMC, and major chemistry journals. This expansion will involve developing criteria for selecting high-quality, recent publications from diverse subfields, thereby ensuring that the dataset reflects both academic rigor and topical relevance.

Improving the semantic verification of QA pairs is also a priority. Current methods, such as fuzzy string matching, may fail to capture nuanced semantic inaccuracies. Future iterations of *ChemRxivQuest* will explore the application of more advanced semantic matching techniques, including transformer-based models and sentence-level embeddings like Sentence-BERT. These approaches are expected to increase the precision of answer validation, thereby reducing the inclusion of subtly incorrect or misleading QA pairs. Comparative benchmarking of different semantic verification strategies will guide the selection of the most robust method for integration.

Finally, the sustainability and long-term impact of *ChemRxivQuest* will be supported through continuous updates and active community involvement. Establishing collaborative platforms—such as GitHub repositories and annotation tools—will enable open contributions, transparent validation workflows, and regular version-controlled releases. This participatory approach will ensure that *ChemRxivQuest* remains a dynamic, community-driven resource, adaptable to evolving research needs in both chemistry and natural language processing.

7 Conclusion

The paper introduces **ChemRxivQuest**, a structured and validated question-answer dataset designed to support natural language processing (NLP) in chemistry. It was developed by extracting and processing content from 155 ChemRxivQuest preprints across 17 chemistry subfields, resulting in 970 carefully verified QA pairs. ChemRxivQuest addresses the lack of chemistry-specific NLP resources and serves as a valuable asset for improving retrieval-augmented systems, fine-tuning large language models, and evaluating scientific knowledge extraction techniques.

The dataset was built using an automated pipeline that integrates optical character recognition (OCR), GPT-based QA generation, and fuzzy-search verification. This pipeline, which is detailed and validated in the paper, can be adapted for constructing domain-specific QA datasets in other scientific fields such as physics or engineering. The authors emphasize the dataset’s high quality, achieved through a similarity-based validation approach that ensures relevance and accuracy—both critical for trustworthy NLP applications.

Despite its strengths, ChemRxivQuest has notable limitations. The authors acknowledge the possibility of hallucinated answers that do not precisely align with source texts, which may require manual review for high-stakes use. They also highlight the bias inherent in relying solely on preprints, which may reflect unverified or preliminary research trends due to the lack of formal peer review.

To address these issues, future plans for ChemRxivQuest include expanding the dataset using peer-reviewed literature, incorporating human expert validation to further ensure accuracy, and regularly updating the dataset to keep pace with new developments in the field. The authors envision a dynamic, community-driven resource that evolves over time.

Overall, ChemRxivQuest stands as a significant contribution to the field of chemistry-focused AI, offering researchers, educators, and industry professionals a reliable and structured source of chemical information. It paves the way for more context-aware NLP systems and sets a benchmark for scientific QA datasets in chemistry.

References

- [1] Daniil A Boiko, Robert MacKnight, Ben Kline, and Gabe Gomes. Autonomous chemical research with large language models. *Nature*, 624(7992):570–578, 2023.
- [2] Jaewoong Choi and Byungju Lee. Accelerating materials language processing with large language models. *Communications Materials*, 5(1):13, 2024.
- [3] Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. Detecting halluci-

- nations in large language models using semantic entropy. *Nature*, 630(8017):625–630, 2024.
- [4] Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yixin Dai, Jiawei Sun, Haofen Wang, and Haofen Wang. Retrieval-augmented generation for large language models: A survey. *arXiv preprint arXiv:2312.10997*, 2:1, 2023.
 - [5] KM Jablonka, P Schwaller, A Ortega-Guerrero, and B Smit. Leveraging large language models for predictive chemistry. *nat. mach. intell.* 6, 161–169, 2024.
 - [6] David M Jessop, Sam E Adams, Egon L Willighagen, Lezan Hawizy, and Peter Murray-Rust. Oscar4: a flexible architecture for chemical text-mining. *Journal of cheminformatics*, 3(1):41, 2011.
 - [7] Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Ye Jin Bang, Andrea Madotto, and Pascale Fung. Survey of hallucination in natural language generation. *ACM computing surveys*, 55(12):1–38, 2023.
 - [8] Qiao Jin, Bhuwan Dhingra, Zhengping Liu, William W Cohen, and Xinghua Lu. Pubmedqa: A dataset for biomedical research question answering. *arXiv preprint arXiv:1909.06146*, 2019.
 - [9] Bonnie Lawlor. Preprints and scholarly communication in chemistry: a look at chemrxiv. *Chemistry International*, 40(4):18–21, 2018.
 - [10] Kyle Lo, Lucy Lu Wang, Mark Neumann, Rodney Kinney, and Dan S Weld. S2orc: The semantic scholar open research corpus. *arXiv preprint arXiv:1911.02782*, 2019.
 - [11] Andres M. Bran, Sam Cox, Oliver Schilter, Carlo Baldassari, Andrew D White, and Philippe Schwaller. Augmenting large language models with chemistry tools. *Nature Machine Intelligence*, 6(5):525–535, 2024.
 - [12] Santiago Miret and Nandan M Krishnan. Are llms ready for real-world materials discovery? *arXiv preprint arXiv:2402.05200*, 2024.
 - [13] Gonzalo Navarro. A guided tour to approximate string matching. *ACM computing surveys (CSUR)*, 33(1):31–88, 2001.
 - [14] Jake Poznanski, Jon Borchardt, Jason Dunkelberger, Regan Huff, Daniel Lin, Aman Rangapur, Christopher Wilhelm, Kyle Lo, and Luca Soldaini. olmocr: Unlocking trillions of tokens in pdfs with vision language models. *arXiv preprint arXiv:2502.18443*, 2025.

- [15] G Appa Rao, G Srinivas, K Venkata Rao, and PVGD Prasad Reddy. A partial ratio and ratio based fuzzy-wuzzy procedure for characteristic mining of mathematical formulas from documents. *IJSC—ICTACT J Soft Comput*, 8(4):1728–1732, 2018.
- [16] Matthew C Swain and Jacqueline M Cole. Chemdataextractor: a toolkit for automated extraction of chemical information from the scientific literature. *Journal of chemical information and modeling*, 56(10):1894–1904, 2016.
- [17] George Tsatsaronis, Georgios Balikas, Prodromos Malakasiotis, Ioannis Partalas, Matthias Zschunke, Michael R Alvers, Dirk Weissenborn, Anastasia Krithara, Sergios Petridis, Dimitris Polychronopoulos, et al. An overview of the bioasq large-scale biomedical semantic indexing and question answering competition. *BMC bioinformatics*, 16:1–28, 2015.
- [18] Johannes Welbl, Nelson F Liu, and Matt Gardner. Crowdsourcing multiple choice science questions. *arXiv preprint arXiv:1707.06209*, 2017.
- [19] Liang Zhang, Katherine Jijo, Spurthi Setty, Eden Chung, Fatima Javid, Natan Vidra, and Tommy Clifford. Enhancing large language model performance to answer questions and extract information more accurately. *arXiv preprint arXiv:2402.01722*, 2024.